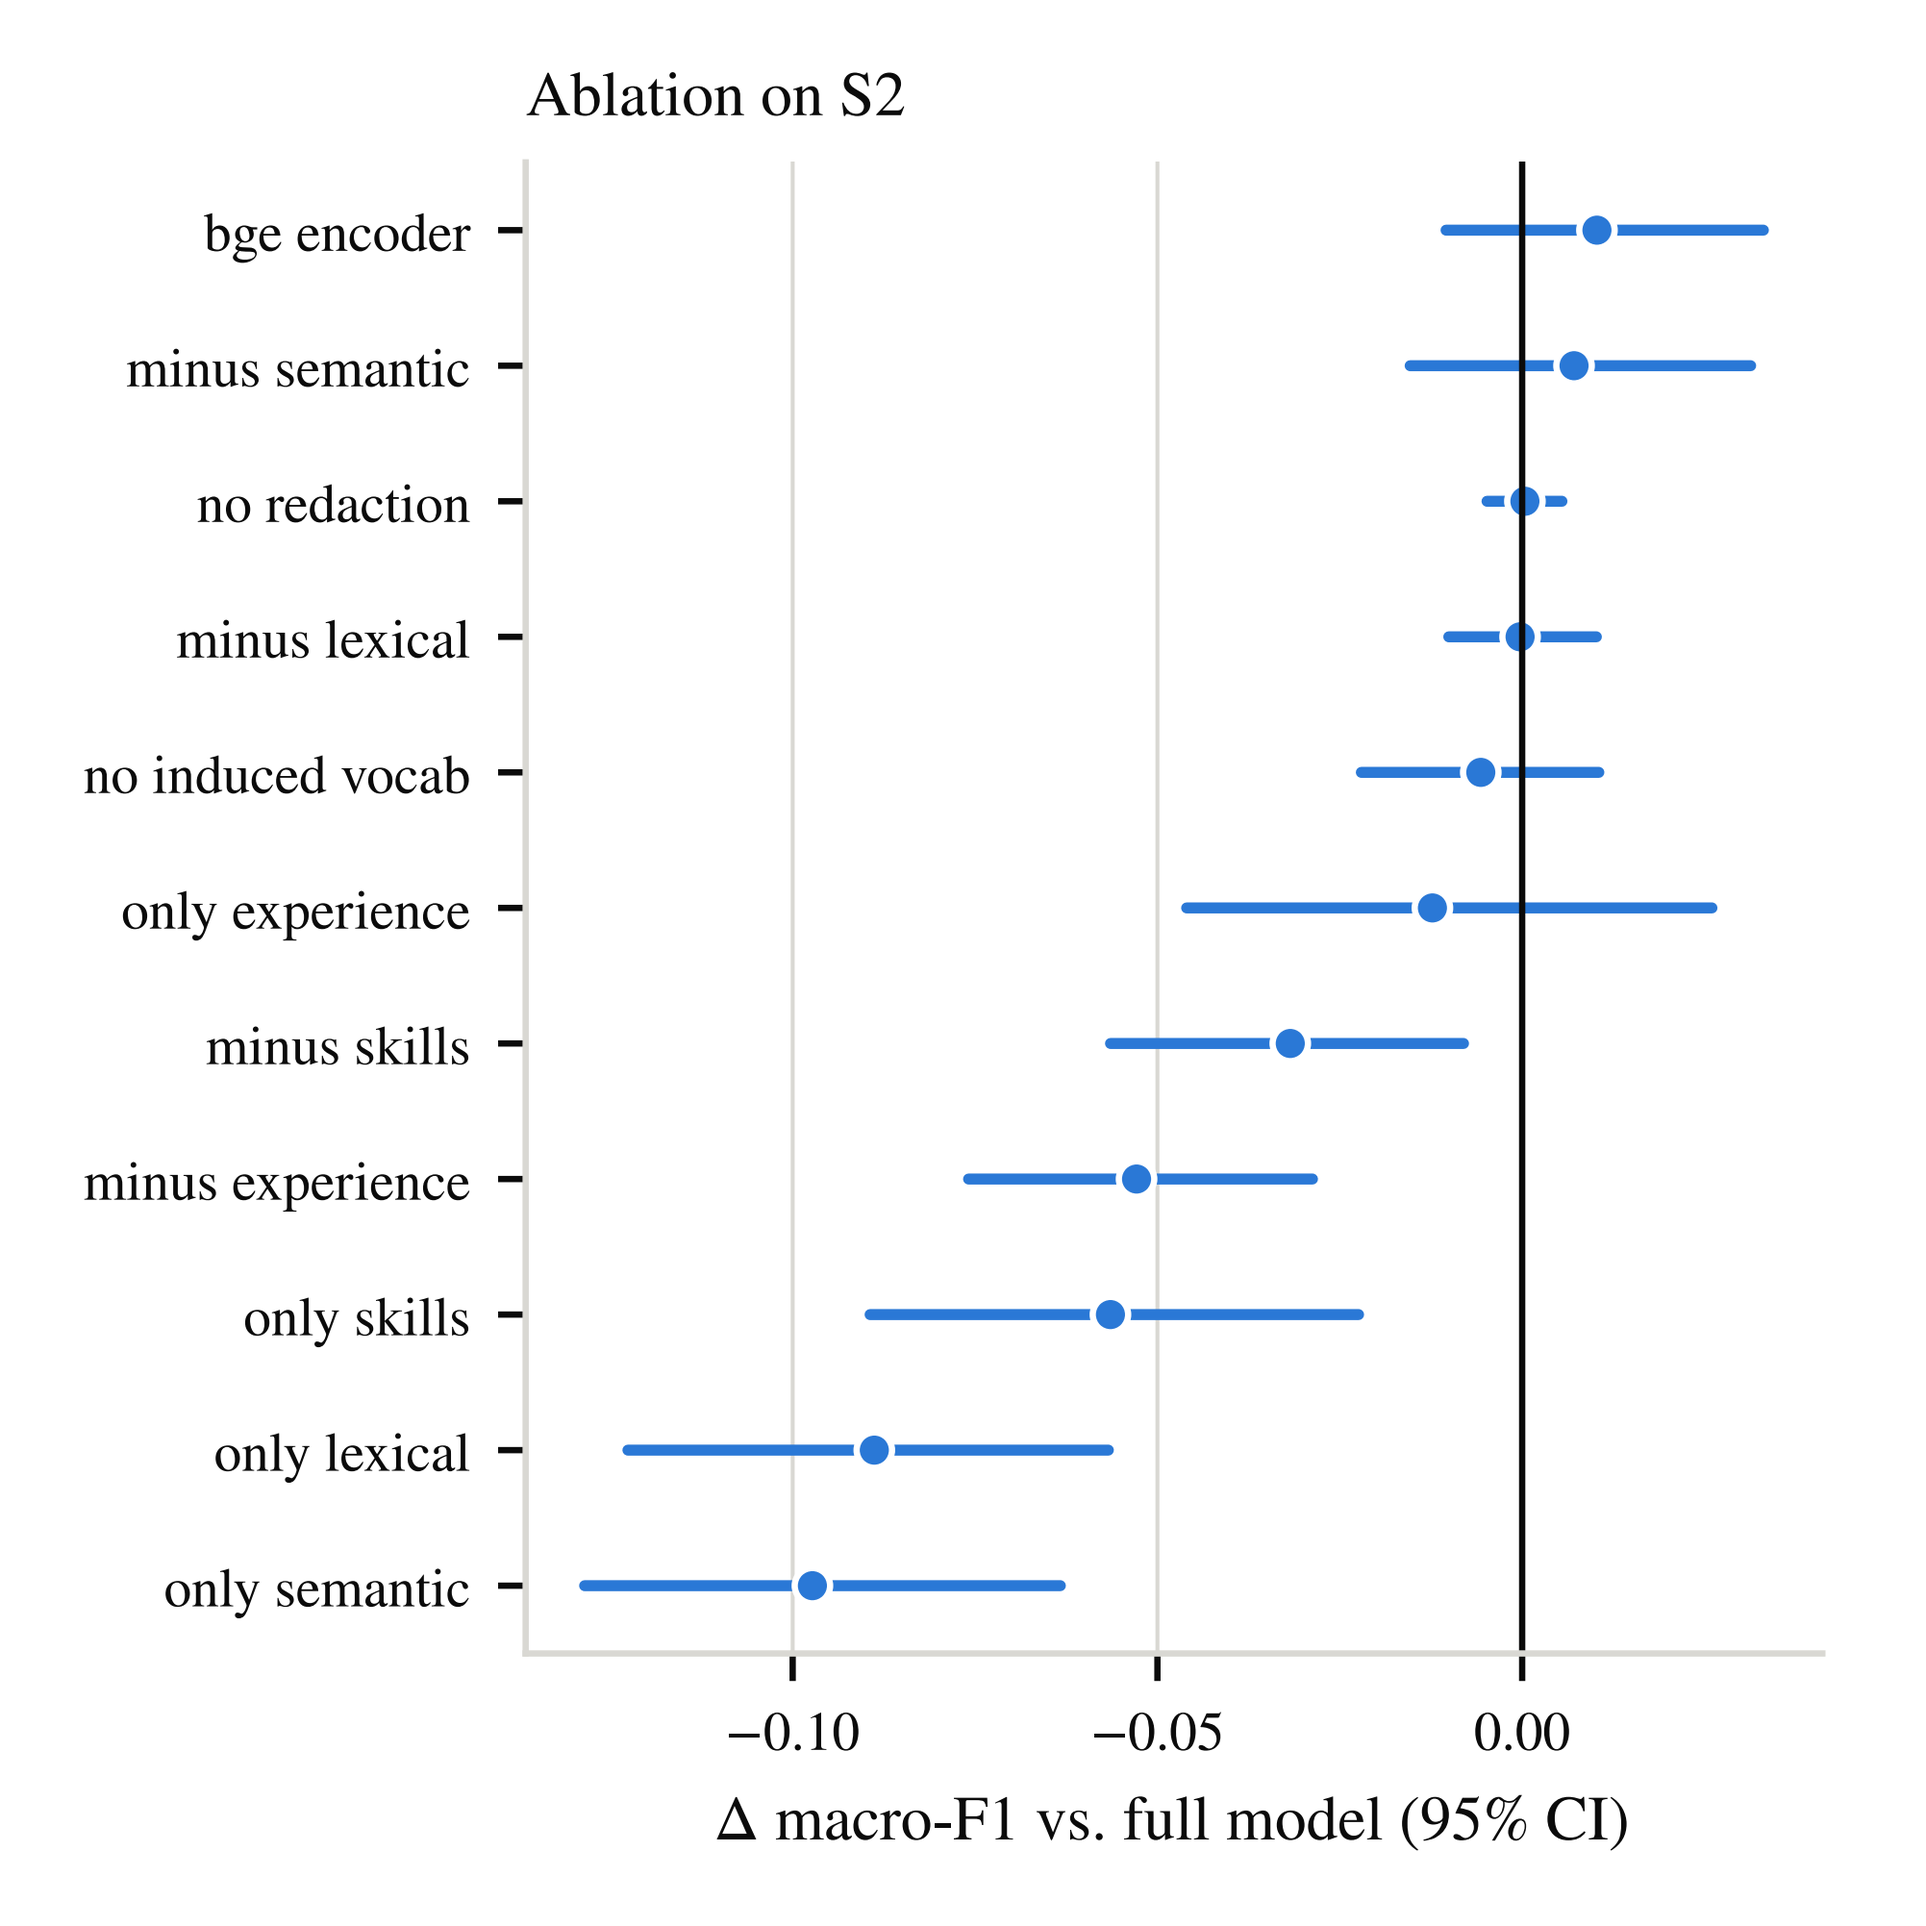
Ablation on S2
bge encoder
minus semantic
no redaction
minus lexical
no induced vocab
only experience
minus skills
minus experience
only skills
only lexical
only semantic
−0.10
−0.05
0.00
Δ macro-F1 vs. full model (95% CI)